

Lab 2: Deep Learning Training Clinic — Robust MLPs + RNNs for Real Sensor Sequences

Projects in ML and AI — Spring 2026 | CSCI-4170/6170

Mini-project (2-3 weeks)

Release Week: Week 6

Lab Time: 90–110 minutes per week

Recommended Duration: 3 weeks

Team Policy: Individual or groups of 2–4

Learning objectives

- Create a reproducible deep learning training pipeline (data loaders, training loop, checkpointing, fixed seeds).
- Diagnose and fix common training failures in real projects: overfitting, unstable loss/NaNs, exploding gradients, and poor generalization.
- Implement and compare a strong MLP baseline (non-sequence) and a sequence model (GRU/LSTM) for windowed sensor data.
- Make model decisions using validation discipline (no peeking at test) and document experiments in a structured log (CSV).
- Produce deployment-readiness artifacts: a saved model checkpoint, a simple inference function, and a basic latency measurement.
- Communicate results clearly via a short report section and a one-page model card (intended use, risks, limitations).

Constraints and allowed methods

Required

- Python + Jupyter Notebook / Google Colab (Pro is fine).
- PyTorch (preferred) or TensorFlow/Keras.
- NumPy, pandas, matplotlib; scikit-learn allowed for metrics only.
- Git repository with clear commit history.
- Experiment Log (CSV): one row per experiment run (model, params, seed, metrics, timestamp).

Allowed (encouraged)

- Early stopping, learning-rate scheduling, dropout, weight decay (L2), gradient clipping.
- Simple imbalance handling (class weights or light re-sampling).
- Mixed precision is allowed only if you can explain what changed and why.

Problem setup: on-device activity recognition from IMU sequences

You are building a prototype for on-device activity recognition using smartphone accelerometer and gyroscope data. This resembles real deployments: windowing and temporal structure, subject generalization, and inference latency constraints.

Dataset (default)

- Use the UCI Human Activity Recognition (HAR) dataset using the "Inertial Signals" (windowed time series).
- You must use sequential windows so an RNN model is meaningful (do not use only engineered summary features).

Leakage and generalization requirement (required)

- Use a subject-disjoint split: no subject IDs shared between train/validation and test.
- Prove it in the notebook by printing the set intersection size (must be 0).
- Your tuning decisions must use only the training subjects (with a validation split). Do not tune on test.

Deliverables (what to submit)

- A Git repo called Lab 2, containing notebooks + any source code.
- Notebook 1: 01_data_and_baselines.ipynb (dataset card + leakage checks + MLP baseline).
- Notebook 2: 02_training_clinic.ipynb (training curves, overfitting control, instability debugging, experiment log).
- Notebook 3: 03_rnn_and_inference.ipynb (GRU/LSTM model, final evaluation, inference function, latency).
- Model Card (1 page, markdown): intended use, evaluation protocol, metrics, limitations, risks.
- Experiment Log (CSV): one row per experiment run (model, params, seed, metrics, timestamp).

What you must build

1. Week 1 Checkpoints

Part A — Dataset card + leakage-safe split + MLP baseline

1. Mini datasheet (short dataset card): motivation, target definition, data source/license, signal description (channels, window length), and limitations/risks.
2. Data sanity checks: confirm shapes (num_windows, T, C), label distribution, and absence of NaN/inf after loading.
3. Leakage audit: implement a subject-disjoint split and print overlap = 0 (required).
4. Baseline 0 (trivial): majority-class predictor or random predictor with fixed seed.
5. Baseline 1 (required): MLP on flattened windows ($T \times C$) with a train/validation split drawn only from training subjects.

- Report one primary metric (accuracy or macro-F1) plus one supporting artifact (confusion matrix and per-class F1 table).

2. Week 2 Checkpoints

Part B — Training clinic: overfitting + instability (Week 2)

- Training curves + reproducibility: plot train/val loss and metric across epochs; fix seeds and record them in your log.
- Overfitting control (required): intentionally induce overfitting (oversized MLP or too many epochs), then fix it using at least two of: dropout, weight decay, early stopping, light noise augmentation, or LR scheduling. Report before/after validation results.
- Instability debug (required): intentionally trigger one failure mode (NaNs from high LR, exploding gradients, or severe underfitting). Then fix it using two appropriate interventions (e.g., lower LR, gradient clipping, improved normalization, adjusted regularization).
- Experiment log discipline: record ≥ 5 runs (teams) or ≥ 3 runs (individual) with key hyperparameters, seed, best validation metric, and a short note per run.

3. Week 3 Checkpoints

Part C — Sequence model + inference readiness (Week 3)

- Core sequence model (required): GRU or LSTM that consumes (batch, T, C) windows and outputs class logits. Include at least one regularization choice (dropout and/or weight decay).
- Training robustness: use gradient clipping (required for RNNs unless you justify why not). Use early stopping or LR scheduling.
- Final evaluation: evaluate your best checkpoint on held-out test subjects. Provide confusion matrix, per-class performance, and a 3–5 sentence error analysis (which classes confuse, why plausible).
- Inference function + latency (required): implement predict_activity(window) \rightarrow (label, probs). Time average inference per window over 100–500 windows on CPU in Colab and interpret feasibility for near-real-time use.
- Model card + limitations: include intended use, non-intended use, evaluation protocol, known failure modes, and a short privacy/ethics note.

Required for Graduate (6170) (choose any ONE)

- Uncertainty: bootstrap confidence intervals for test macro-F1 (≥ 500 resamples).
- Interpretability: saliency/gradient-based attribution over timesteps/channels with a short discussion of a failure mode.
- Generalization: report performance variability across subjects and discuss deployment risk.

Weekly milestones and Thursday check-ins

Week	Milestone (due by Thursday lab session)	Evidence to show in the 3–5 minute explanation
Week 1	Datasheet + subject-safe split + MLP baseline	Notebook outputs; overlap=0 proof; baseline metrics; confusion matrix; commits
Week 2	Training clinic: overfitting fix + instability fix + experiment log	Before/after curves; one debug story; experiment log rows; commits
Week 3	GRU/LSTM + test eval + inference + latency + model card	Final metrics; confusion matrix; inference demo; timing results; model card; commits

End-of-lab explanation (required for full credit)

At the end of each Thursday lab session, your team must give a 3–5-minute explanation using this structure:

- What changed since last session? (show git commits or notebook diffs)
- What did you measure and what did you learn? (show one table/plot)
- One technical decision: what you tried, expected outcome, actual outcome, next step
- Instructor verification questions (any team member may answer)

Note: If the team cannot demonstrate evidence of progress and shared understanding, the oral-check component may be reduced for that week.

Submission checklist

- All notebooks run top-to-bottom without manual edits; all random seeds are fixed and reported.
- Repo contains an Experiment Log CSV with ≥ 5 runs (teams) or ≥ 3 runs (individual).
- Repo contains: baseline results, training-clinic evidence (curves + fixes), RNN results, and final test evaluation.
- Model Card included as markdown in the repo.
- Inference function works and latency measurement is reported.